



**TRIBHUVAN UNIVERSITY
INSTITUTE OF ENGINEERING
THAPATHALI CAMPUS**

**A MINOR PROJECT REPORT
ON
HUMAN PRESENCE DETECTION USING DISTRIBUTED INFERENCE ON
RESOURCE CONSTRAINED EDGE DEVICES**

Submitted By:

Ishwor Raj Pokhrel (THA079BEI013)
Pramit Khatri (THA079BEI020)
Preeti Rijal (THA079BEI028)
Sampada Ghimire (THA079BEI036)

Submitted To:

Department of Electronics and Computer Engineering
Thapathali Campus
Kathmandu, Nepal

In partial fulfillment for the award of the
Bachelors Degree in Electronics, Communication and Information Engineering

Under the Supervision of
Asst. Prof. Suwarna Lingden

May , 2026



**TRIBHUVAN UNIVERSITY
INSTITUTE OF ENGINEERING
THAPATHALI CAMPUS**

**A MINOR PROJECT REPORT
ON
HUMAN PRESENCE DETECTION USING DISTRIBUTED INFERENCE ON
RESOURCE CONSTRAINED EDGE DEVICES**

Submitted By:

Ishwor Raj Pokhrel	(THA079BEI013)
Pramit Khatri	(THA079BEI020)
Preeti Rijal	(THA079BEI028)
Sampada Ghimire	(THA079BEI036)

Submitted To:

Department of Electronics and Computer Engineering
Thapathali Campus
Kathmandu, Nepal

May , 2026

CERTIFICATE OF APPROVAL

The undersigned certify that they have read and recommended to the **Department of Electronics and Computer Engineering, IOE, Thapathali Campus**, a minor project work entitled “**Human Presence Detection Using Distributed Inference On Resource Constrained Edge Devices**” submitted by **Ishwor Raj Pokhrel, Pramit Khatri, Preeti Rijal, Sampada Ghimire** in partial fulfillment for the award of Bachelor of Engineering in Electronics, Communication and Information Engineering. The project was carried out under the special supervision and within the time prescribed by the syllabus.

We found the students to be hardworking, skilled and ready to undertake any related work to their field of study and hence we recommend the award of partial fulfillment of Bachelor’s Degree of Electronics, Communication, and Information Engineering.

Project Supervisor
Asst. Prof. Suwarna Lingden
Department of Electronics and Computer
Engineering
Thapathali Campus, IOE

External Examiner
Er. Birodh Rijal
Co-Founder and CEO, Integrated ICT Pvt.
Ltd.
Jwagal, Kupondole, Lalitpur, Nepal

Head of Department
Asst. Prof. Umesh Kanta Ghimire
Department of Electronics and Computer
Engineering
Thapathali Campus, IOE

May , 2026

DECLARATION

We hereby declare that the project entitled, **“Human Presence Detection Using Distributed Inference On Resource Constrained Edge Devices”** for the partial fulfillment of the requirements for the award of the degree of Bachelor of Engineering in **Electronics, Communication and Information Engineering** is a bonafide report of the work carried out by us. These materials contained within the report have not been submitted to any University or Institution for the award of any degree, and we are the only author of this complete work and no sources other than the listed ones have been used in this work.

Ishwor Raj Pokhrel (THA079BEI013)

Pramit Khatri (THA079BEI020)

Preeti Rijal (THA079BEI028)

Sampada Ghimire (THA079BEI036)

Date: May , 2026

COPYRIGHT

The authors has agreed that the library, Department of Electronics and Computer Engineering, Thapathali Campus, may make this report freely available for inspection. Moreover, the authors have agreed that the permission for extensive copying of this project work for scholarly purpose may be granted by the professor/lecturer, who supervised the project work recorded here in or, in their absence, by the Head of the Department. It is understood that the recognition will be given to the authors of this report and to the Department of Electronics and Computer Engineering, IOE, Thapathali Campus in any use of the material of this report. Copying of publication or other use of this report for financial gain without approval of the Department of Electronics and Computer Engineering, IOE, Thapathali Campus and authors' written permission is prohibited.

Request for permission to copy or to make any use of the material in this project in whole or part should be addressed to the Department of Electronics and Computer Engineering, IOE, Thapathali Campus.

ACKNOWLEDGEMENT

We sincerely wish to thank our project supervisor, **Asst. Prof. Suwarna Lingden**, for his guidance, helpful advice, and continued support in the project.

We sincerely thank **Asst. Prof. Umesh Kanta Ghimire**, Head of the Department, and the whole Department of Electronics and Computer Engineering for providing us with the required facilities and a conducive learning atmosphere.

We thank all our faculty members, friends, and all the other individuals, both directly or indirectly involved in this project for giving us their valuable suggestions and supportive encouragement.

Ishwor Raj Pokhrel	(THA079BEI013)
Pramit Khatri	(THA079BEI020)
Preeti Rijal	(THA079BEI028)
Sampada Ghimire	(THA079BEI036)

ABSTRACT

The field of monitoring and surveillance requires the use of high resolution image sensors but the computational power required to process such data is very high and becomes a significant challenge for edge device processing. In order to process the high resolution image in an edge device requires an aggressive spatial downsampling which destroys the fine details necessary for reliable detection of small and distant subjects. This work presents a distributed inference architecture for real-time Human Presence Detection (HPD) on microcontroller based systems. Rather than downsampling the full frame, the pipeline partitions each captured image into a grid of smaller tiles, allowing each tile to be inferred independently at full local resolution. The resulting tile-level predictions are passed to a lightweight aggregator, which combines outputs across all tiles to produce a single classification score. End-to-end benchmarking reveals that frame capturing and tiling completes in 4 ms, aggregator classification in 0.342 ms, and takes 472 ms for feature transmission per tile. The dominant bottleneck is tile level inference on the ESP32-S3, which requires 2,657 ms per tile. The total critical-path latency comes to 18.9 seconds approximately against the 1,000 ms target. The results side with the architectural correctness for all architecture stages. They also highlight two issues that need to be addressed: Memory bound inference latency and Sequential tile fragmentation . These are the areas that need optimisation in future work.

Keywords: *Human presence detection, distributed tile-based inference, local binary convolutional neural network, attention mechanism, TinyML, post-training quantization, ESP32, edge computing.*

TABLE OF CONTENTS

CERTIFICATE OF APPROVAL	i
DECLARATION	ii
COPYRIGHT	iii
ACKNOWLEDGEMENT	iv
ABSTRACT	v
LIST OF FIGURES	vii
LIST OF TABLES	viii
LIST OF ABBREVIATIONS	viii
1. Introduction	1
1.1. Background	2
1.2. Motivation	3
1.3. Problem Definition	4
1.4. Problem Definition	4
1.5. Objectives	6
1.6. Scope and Applications	6
2. LITERATURE REVIEW	8
3. REQUIREMENTS ANALYSIS	10
3.1. Hardware Requirements	10
3.2. Software & Framework Requirements	10
3.3. Non-functional Requirements	10
4. SYSTEM ARCHITECTURE AND METHODOLOGY	12
4.1. Model Selection and Optimization Pipeline	12
4.2. Hardware Requirements and system configuration	19
4.3. Hardware Used	24
4.4. Evaluation Metrics and Performance Framework	27
5. RESULTS AND ANALYSIS	28
5.1. MobileNetV2 Architecture Performance	28
5.2. MobileNetV2 Test-Set Evaluation	29
5.3. LBCNN Architecture Performance	29

5.4. Test-Set Evaluation	31
5.5. Comparative Analysis: LBCNN vs. MobileNetV2	31
5.6. Overall Comparative Analysis	32
5.7. System Overview and Evaluation Methodology	33
6. CONCLUSION	39
6.1. Achievements	39
6.2. Limitations	39
6.3. Future Work	40
6.4. Final Remarks	42
A. APPENDICES	43
A.1. PROJECT BUDGET	43
A.2. PROJECT TIMELINE	43

LIST OF FIGURES

Figure 4-1. Multi-Tile Attention-Based CNN Architecture	13
Figure 4-2. Detailed Schematic of the Local Binary Convolution Module	14
Figure 4-3. System Architecture showing main components and their interactions.	20
Figure 4-4. Detailed Data Flow Diagram of the Tiling and Inference Pipeline. . .	21
Figure 4-5. Wiring Interface between Programming (FTDI), Processing (ESP32-S3), and Visualization (I2C Liquid Crystal Display (LCD)) Units. . . .	24
Figure 5-1. MobileNetV2 training dynamics and test-set evaluation results. . . .	28
Figure 5-2. LBCNN training dynamics and test-set evaluation results.	30

LIST OF TABLES

Table 4-1. Wiring the FTDI to ESP32-CAM	25
Table 4-2. Wiring the I2C LCD to ESP32	26
Table 5-1. MobileNetV2 Training Run Summary	29
Table 5-2. Per-Class Classification Report on the Test Set ($n = 2432$) for MobileNetV2	29
Table 5-3. Training Run Summary	30
Table 5-4. Per-Class Classification Report on the Test Set ($n = 92$)	31
Table 5-5. Performance Comparison: LBCNN vs. MobileNetV2	31
Table 5-6. End-to-End Pipeline Latency Breakdown	33
Table 5-7. Worker Node Per-Tile Inference Latency	35
Table 5-8. Aggregator Node Processing Latency	36
Table 5-9. Critical Path Latency vs. Target Frame Budget	37
Table 5-10. Model Quantisation Summary	37
Table A-1. Budget Estimation	43
Table A-2. Gantt Chart for Project Timeline	43

1. INTRODUCTION

Convolutional Neural Network (CNN) perform well at detecting objects in controlled environments. But moving to the real world, high-resolution processing is still a challenge when processing on edge system. Most small processors can't store and process large images directly. A common way to handle this is to shrink the image before processing. This causes the loss of smaller details that are valuable for accurate detection of objects. This project is heavily influenced from field of precision agriculture. Where, large areas are analysed by breaking them into smaller sections and processing each one separately. Here, We applied the same idea, combined with distributed edge computing system in order to build a prominent solution for real-time processing Human Presence Detection (HPD). We chose HPD as our focus to solve a specific problem as Detecting people in large, open spaces requires fine visual details. Traditional downsampling destroys that granularity. Our approach splits a high resolution image into smaller square tiles then each tile will be processed by a microcontroller spreading the workload across many similar devices instead of relying on a single powerful one. As a result, we preserve the image quality needed to detect people reliably while processing on lower power devices. This method also has practical benefits beyond accuracy as all processing happens locally on the device without needing the internet or the cloud compute. It reduces latency and preserves privacy.

1.1. Background

Common object detection models like You Only Look Once (YOLO) and Faster architecture like R-CNN are designed to work with small input images with sizes like 416×416 or 640×640 pixels. However, cameras used in surveillance or agricultural automation can capture images as large as 4000×3000 pixels which is nearly ten times larger than what these models expect.

The Problem of Vanishing Detail

Surveillance systems often capture very large images but most CNN models only accept small, fixed size inputs. To bridge this gap, the image is shrunk down before being fed into the model which is known as the Spatial Information Attrition problem which causes serious data loss. Consider a frame that is 4000 pixels wide being reduced to 400 pixels. That's a 90% reduction in width and 99% of the total pixel data is thrown away.

Small objects suffer the most from this process. Take a person standing at a distance. They might occupy a 40×80 pixel region in the original image but after downsampling, that same person is reduced to just 4×8 pixels which is roughly the size of a thumbnail on a postage stamp

At that size, a CNN cannot do its job properly. Every CNN has a minimum region it needs to see in order to pick up useful features which is called the Effective Receptive Field (ERF) [?]. A 4x8 pixel blob falls well below that minimum which does not have enough details to identify limbs, a torso, or any other feature that distinguishes a person from background noise. The model ends up trying to classify a blurry smudge, which can look surprisingly identical to a shadow or a bush.

Lessons from High-Precision Agriculture

In precision agriculture, satellites and drones are used to capture very high resolution image of fields. Researchers then analyse these images to spot individual plants, plots, pests, or broken irrigation pipes. These objects are extremely tiny and often take up less than 0.1% of the total image [?]. This is known as the Small Object Detection Small Object Detection (SOD) problem.

To handle this, agricultural models use a technique called Tiling, also known as Patch based Inference [?], [?]. Instead of shrinking the whole image, the image is cut into a grid of smaller sections called tiles then Each tile is processed separately which preserves what is called the Ground Sampling Distance Ground Sampling Distance (GSD). In simple terms, GSD refers to how much real world area a single pixel represents. By keeping the tiles at full resolution, no spatial detail is lost. One well-known implementation is the Slicing Aided Hyper Inference Slicing Aided Hyper Inference (SAHI) framework [?]. Studies using SAHI show that tiling can improve small object detection accuracy by 20 to 30% compared to simply resizing the full image [3].

However, these agricultural solutions have a limitation which is; they are built for powerful workstations and are typically run offline and were not designed for real time processing on small, low powered devices.

Growing AI and Evolving Edge Computing

Artificial Intelligence is no longer confined to large data centres in the cloud rather is increasingly happening closer to where data is being collected. This shift is known as edge computing. By 2025, it is estimated that over 75% of enterprise data is created and processed outside traditional data centres [?]. This change is being driven by the need for speed. Applications like autonomous surveillance and industrial monitoring cannot afford the delay of sending data to a remote server and waiting for a response which can be solved by processing locally [?].

Edge computing works through a network of small, devices ranging from simple IoT components to more sophisticated hardware like NVIDIA Jetson where each device

performs some or all of the AI processing locally, without relying on the internet and cloud.

Modern AI models are growing more large and complex and single edge device often does not have enough memory along with processing power to run them. This has led to a new dimension namely Distributed Edge Inference [?] where Instead of running the entire model on one single device, the workload is split across several interconnected devices where each device is responsible for smaller portion of larger workload.

1.2. Motivation

««««< HEAD Modern cameras are getting better and better as 4K and 8K sensors are now common in security and industrial monitoring. But AI models have not kept up. Most high performance models still expect small, fixed size inputs. This gap is the core motivation for this research. Currently, there are two common ways to deal with this mismatch but neither works well. The first option is to shrink the image size. This is rapid and cheap, but it removes fine details. Small or distant subjects become impossible to detect reliably. The second option is to use more powerful hardware. gpu can handle large images, but they consume a lot of power. This makes them unsuitable for remote locations or battery powered devices. So some middle ground is needed. his research aims to build an architecture that keeps the full detail of the original image, while distributing the processing load across multiple small, affordable devices.

1.3. Problem Definition

This research addresses a core problem in high resolution computer vision. As cameras get better and better, the gap between how much data they capture and speed of processing keeps growing. This is sometimes called the Information Computation Paradox. It breaks down into three specific problems. **The Scaling Bottleneck and Receptive Field Constraints:**Most CNN models like ResNet or You Only Look Once (YOLO) expect a fixed input size, typically 416×416 or 640×640 pixels [?] . When a 4K video frame is forced into that size, the image must be heavily resized. This resizing acts like a filter by smoothing out fine details and removes the sharp edges that define small objects. The result is that the region a CNN focuses on becomes larger than the object it is trying to find. The network fails to detect it [?].

Hardware and Memory Constraints:Edge devices are small and have limited memory so running a large AI model directly on one while also holding a high resolution image in memory is often not possible. The device either crashes with an Out of Memory (OOM) error or slows down due to overheating. There are currently no widely adopted frameworks that allow a single inference task to be split spatially across a group of low

power devices [?]. ===== Security and industrial monitoring now frequently use 4k and 8k sensors. As Cameras are getting better and better. But AI models haven't kept pace. Most high performance models still expect small, fixed sized inputs which is the motivation for this work. There are two common ways to deal with this mismatch and neither works well. The first option is to reduce the image size. It is inexpensive and fast, but loses fine detail. Small or distant objects cannot be reliably detected with this method. The other option is to use more powerful hardware. The Graphics Processing Unit (GPU) can process very big images but they use a lot of compute power. This makes them not suitable for battery operated devices or at remote locations. So we need some kind of middle ground that address the issues mentioned above. The goal of his research is to develop an architecture that keeps the detail of the original image, but distributes the processing load among many small and cheap devices.

1.4. Problem Definition

The research works on the fundamental problem in high res image processing. As camera are getting better and better the amount of pixels they capture and rate of processing is getting larger and larger. Which is popularly known as information computation paradox. It further breaks down into 3 particular problem :

The Scaling Bottleneck and Receptive Field Constraints:

Most CNN models such as ResNet or YOLO are designed to take a fixed input size such as 416×416 or 640×640 pixels [?] . That size is going to require massive resizing of the image to fit a 4K video frame. This resizing is essentially a filter that blurs fine detail and removes the hard edges that make small objects pop. As a consequence, the region a CNN focuses on becomes larger than the object it is trying to find. The network cannot recognize it [?].

Hardware and Memory Constraints: Edge devices are small with limited memory, so it is often not possible to run a large AI model directly on one and hold a high resolution image in memory. The device hangs with an Out of Memory (OOM) error or starts lagging due to heat. Currently, there are no common frameworks available that would allow a single inference task to be partitioned in space among a cluster of low power devices [?]. >>>> a0f91c8 (humanized section 1 2 3)

The Aggregation Challenge: Dividing the image into smaller tiles create a problem near the boundaries of each tile. If a person is split among multiple tiles then no single device will have complete information to deduce if the object is human or not. The major challenge comes to splitting the image and designing the aggregator algorithm in such

a way that it can still make decision based on partial decision from each node without using a much powerful hardware. [?]

1.5. Objectives

- To build a smaller light weight model that can process a higher resolution image by dividing it into smaller tiles.
- To detect the presence of human by utilizing the above model on a distributed architecture of edge devices.

1.6. Scope and Applications

This project address the two different dimensions, distributed computing and computer vision. Here Human Presence Detection is taken as a main test case but the model and the system architecture itself is not limited to this particular task only. By making the high resolution image processing more affordable and less complex to process on edge devices this project targets to solve problems across more than one fields.

««««< HEAD

- **High-Precision Industrial Safety:** Large factory floors need constant monitoring. Small but dangerous events, such as a hand getting too close to moving machinery, must be caught early. This system can detect such events at high resolution without the need for expensive industrial computers.
- **Privacy-Preserving Smart Homes:** Because each tile is processed locally and only the final results are shared, raw images never leave the device. This makes the system a good fit for home presence sensing and elderly care, where privacy is a major priority.
- **Intelligent Energy Management of Heating, Ventilation, and Air Conditioning (HVAC) Systems:** Large offices and commercial buildings need accurate occupancy data to manage their Heating, Ventilation, and Air Conditioning HVAC systems efficiently. Traditional Passive Infrared (PIR) sensors struggle to detect people who are sitting still to solve this a distributed CNN based system can detect occupants regardless of whether they are moving or stationary. This information can be fed directly into a Building Management System (BMS), allowing the HVAC system to adjust airflow and temperature room by room in real time.

=====

- **High-Precision Industrial Safety:** Large industries often need surveillance for various issues like hazards that impact machineries and human life. For example detecting and warning the worker if hand goes in dangerous zone of moving machine.
- **Privacy-Preserving Smart Homes:** As each image is processed in distributed architecture without needing any connection and data transfer over cloud this can be used in smart homes without worrying about the data privacy risks.

- **Intelligent Energy Management of HVAC Systems:** Large commercial buildings use Heating, Ventilation, and Air Conditioning HVAC systems which are power hungry and require effective management. In order to effectively manage the HVAC systems we need the occupancy data i.e. If a human is present in a room or not. Traditionally PIR sensors are being used but they struggle to detect stationary humans. But an image processing system can easily detect presence of humans regardless of motion or not. This information can be used to manage HVAC systems. »»»> a0f91c8 (humanized section 1 2 3)

This becomes effective for static or slow moving objects. The system works best in scenes where changes happen slowly enough that the processing delay does not cause missed detections.

2. LITERATURE REVIEW

Standard sensors struggle to detect people who are not moving. Ahmad et al. [?] studied hybrid systems that combine PIR and Radar sensors to improve detection in smart buildings. Combining the two reduced false alarms, the system still showed nearly 0% reliability when occupants were completely still. The study did not explore visual confirmation as a solution. This leaves a gap that computer vision can fill, by detecting the small physical movements that people make even at rest.

Running vision models on small hardware requires significantly reducing input size. Nagel et al. [?] examined INT8 Post-Training Quantisation, a technique that converts model weights from 32-bit floating point to 8-bit integers. This reduces memory usage by four times with less than 1% reduction in accuracy. However, the study does not address the specific constraints of microcontroller based devices like the ESP32-S3, which has only 512 KB of SRAM [?] . Deploying on such hardware requires additional instruction-level optimisations that the paper does not cover.

Wu et al. [?] explored further compression through the Q-CNN framework. Rather than simply reducing weight size, Q-CNN focuses on matching the output of compressed layers to the original model as closely as possible. This allows models to be compressed by up to 20 times while maintaining strong classification performance. The limitation is that the study only tests single-device deployment and does not address how to process high-resolution images across multiple low-power nodes.

Iandola et al. [?] introduced SqueezeNet, which replaces large convolutional filters with smaller 1×1 layers called Fire Modules. The result is a model that matches AlexNet accuracy using 50 times fewer parameters making it well suited for mobile and embedded devices. However, SqueezeNet was not designed for ultra-low-bit operations, which are needed when deploying on the most power-constrained hardware used in this project.

Juefei-Xu et al. [?] took efficiency further with Local Binary Convolutional Neural Networks Local Binary Convolutional Neural Network (LBCNN). Instead of learned filters, LBCNN uses fixed binary kernels that do n't change during training process. This reduces the number of parameters by 169 times and speeds up inference in the same way. The drawback is that LBCNN has no builtin way to handle very large resolution images, which are too large to fit in the memory of most microcontrollers.

Li et al. [?] studied how large images can be broken into grids of smaller sections for processing in deep learning systems. Their findings confirmed that tiling does preserves fine spatial detail that is lost when an image is resized. The study was conducted on

GPU clusters, however, and does not provide a communication protocol suited for small, low-power device clusters

Akyon et al. [?] addressed small object detection directly through the SAHI framework which converts a high-resolution image into overlapping patches and runs inference on each one separately. This showed that accuracy of detection can be improved by 20 to 30% [?]. Even though the framework was designed for powerful workstations and is computationally heavy and was not designed for real time application the founding principles can still be applied.

Jeong et al. [?] studied low latency communication protocols for network of IoT devices. Results from the study showed that peer-to-peer protocols work best with latency as low as 1.115ms which is fast enough for our real time synchronisation. Even though study was not specifically done on image data transmission which requires far more data transfer than simple sensors we can still assume similar outcome. This project addresses all of these gaps together by deploying a real-time distributed tiling system on ESP32-S3 microcontroller based architecture, combining spatial decomposition, model quantisation, and low-latency communication in a single pipeline designed for resource-constrained hardware.

3. REQUIREMENTS ANALYSIS

The development of a distributed, tile-based HPD system requires a multi level analysis of hardware constraints, software dependencies, and network reliability. The section below discusses the basic requirements needed for the project to be feasible.

3.1. Hardware Requirements

Given the ESP32-S3 chip in the design, the choice of hardware is based on the speed up and power consumption needs related to Artificial Intelligence (AI).

- **Camera Module:** A high-resolution camera module(OV2640) capable of capturing images at higher resolution than traditional image processing systems, which will then be partitioned.
- **Input Node:** A master node responsible for taking images from the image sensor; partitioning them into multiple tiles and distributing them to edge nodes.
- **Edge Computing Nodes:** Multiple ESP32-S3 Microcontroller Units (MCUs) are required. These must support the *XtensaLX7* vector instructions to handle the tile-level CNN inference.
- **Aggregator Node:** A simple ESP32 responsible for collecting feature vectors and executing the final classification layer.
- **Communication Interface:** A low-latency communication system for inter-node communication Espressif Non-Beaconing Wireless (ESP-NOWs).

3.2. Software & Framework Requirements

To achieve ultra-low bit quantization and distributed processing, the following software stack is required:

- **Development Environment:** Arduino IDE using C/C++.
- **AI Inference Library:** The Espressif Deep Learning library (ESP-DL) makes use of efficient kernels for quantization 8-bit or 16-bit computation.
- **Model Optimization:** Initial training and conversion of models with TensorFlow Lite for Microcontrollers (TFML) are followed by Quantization-Aware Training (QAT).
- **Tiling Logic:** A custom written algorithm to handle the tiling of image and memory management for each node processing the tile.

3.3. Non-functional Requirements

- **Latency:** Low latency is required for a real time system.

- **Scalability:** The number of processing node should be adjustable instead of rigid as per level of resolution and required Frames Per Second (FPS).
- **Robustness:** Failure of a single intermediate processing node should not impact the processing and it should adjust to the available working nodes.

4. SYSTEM ARCHITECTURE AND METHODOLOGY

4.1. Model Selection and Optimization Pipeline

4.1.1. MobileNetV2 Architecture

The proposed architecture uses a CNN model based on mobilenetv2 with optimization for binary classification, which is expressed mathematically $f_{\theta} : \mathbb{R}^{N \times 1 \times H \times W} \rightarrow \mathbb{R}^{N \times 2}$. The design incorporates a compromise between deep feature extraction and MCU hardware limitations.

Grayscale Adaptation and Transfer Learning

To reduce memory usage and improve computation time, the model is adjusted to handle single-channel images. As the weights of the model were initially trained on RGB images, the first convolutional layer weights ($W \in \mathbb{R}^{C_{out} \times 3 \times k \times k}$) are averaged across the channel dimension:

$$Y = 0.299R + 0.587G + 0.114B \quad (4-1)$$

This technique retains low level feature detection, such as edge detectors and textures, while allowing the model to function on simplified Static Random Access Memory Static Random-Access Memory (SRAM) efficient inputs.

Classifier Head and Loss Function

The final classification is performed by a narrow bottleneck head ($1280 \rightarrow 64 \rightarrow 2$) using ReLU activation and dropout (0.2 and 0.1) to prevent overfitting. The resulting logits $\hat{y}$ are evaluated using Cross-Entropy Loss:

$$L = - \sum_{i=1}^2 y_i \log \left(\frac{e^{\hat{y}_i}}{\sum_j e^{\hat{y}_j}} \right) \quad (4-2)$$

4.1.2. Local Binary Convolution Based Architecture




Figure 4-1. Multi-Tile Attention-Based CNN Architecture




Figure 4-2. Detailed Schematic of the Local Binary Convolution Module

LBCLayer Overview

The LBCLayer is a hybrid convolutional layer designed to combine fixed binary convolutions (anchors) for efficient sparse structural feature capture with a learnable 1×1

convolution for task-specific adaptation. It is particularly suitable for resource-efficient networks and sparse feature representations, inspired by research on Local Binary Convolutions (LBC) and sparse binary kernels.

Mathematically, for an input $x \in \mathbb{R}^{N \times C_{in} \times H \times W}$, the output is defined as:

$$x_{out} = \text{Conv1x1}(\text{ReLU}(\text{Conv}_{\text{binary}}(x))) \quad (4-3)$$

Where:

- $\text{Conv}_{\text{binary}}$ uses fixed ternary weights $\in \{-1, 0, 1\}$.
- Conv1x1 is trainable and combines channels efficiently.

LBCLayer Components

1. Fixed Anchor Convolution

The purpose of this component is to extract general features without learning, thereby reducing the parameter count.

- **Weight generation:** The first step is the creation of weights matrix $W \sim \mathcal{N}(0, 1)$ and its sparsifying with parameter s :

$$\text{binary_weight}_{i,j,k,l} = \begin{cases} 1 & \text{if } W_{i,j,k,l} > \text{threshold} \\ -1 & \text{if } W_{i,j,k,l} < -\text{threshold} \\ 0 & \text{otherwise} \end{cases}$$

- The sparsity parameter s will define how many zeros will there be in the matrix.
- This matrix of weights is saved as a buffer (`register_buffer`), which means it won't be trained afterwards.

The main idea is to capture general directional/edge patterns efficiently similar to handcrafted filters like Sobel, but stochastically generated.

2. Non-linearity (ReLU)

After the fixed binary convolution, non linearity is applied by ReLU: $x' = \text{ReLU}(x)$. Non linearity is introduced to help approximate complex functions, which also maintain sparsity and improving computational efficiency since there will be a lot of zeros.

3. Learnable 1x1 Convolution

The 1×1 kernel combines sparse feature maps into task relevant channels while keeping computational costs low.

$$x_{\text{final}} = W_{1 \times 1} * x' \quad (4-4)$$

Only this part is trainable, allowing the layer to adapt to specific tasks like human detection.

Forward Pass Flow

For a single LBCLayer, the flow is as follows:

1. The input to the layer is $x \in \mathbb{R}^{N \times C_{\text{in}} \times H \times W}$.
2. Perform binary convolution using fixed parameters: $x_1 = \text{Conv}_{\text{binary}}(x)$.
3. Use ReLU activation function: $x_2 = \text{ReLU}(x_1)$.
4. Perform 1×1 convolution using learnable parameters: $x_{\text{out}} = \text{Conv}_{1 \times 1}(x_2)$.
5. Output: $x_{\text{out}} \in \mathbb{R}^{N \times C_{\text{out}} \times H \times W}$.

Advantages of LBCLayer

- **Computational efficiency:** Fixed sparse weights reduce multiply adds, and 1×1 learnable convolutions are computationally cheap.
- **Stability:** Binary anchors are not updated, which avoids noisy gradients and the network only optimizes the final 1×1 convolution.
- **Sparsity Control:** The sparsity parameter allows control over network complexity and feature density so that encouraging sparse representations that improve generalization.
- **Compatibility with Tilings:** LBCLayer is very well suited for tiling based systems such as Human detector where individual tiles can be processed separately before being assembled.
- **Biological Inspiration:** Sparse and binary inspired filters mimic biological visual cortex processing that is helpful for edge detection and texture recognition.

Mathematical Summary

Considering the inputs x , the summary of the operation and gradients are as follows:

$$\text{Anchor weight: } A = f_{\text{binary}}(W, s)$$

$$x_1 = x * A$$

$$x_2 = \text{ReLU}(x_1)$$

$$x_{\text{out}} = x_2 * W_{1 \times 1}$$

Gradients: $\frac{\partial L}{\partial W_{1 \times 1}} \neq 0$, whereas $\frac{\partial L}{\partial A} = 0$. In this case, $*$ denotes a convolution, s denotes the sparsity threshold, and A denotes the fixed binary anchor weights.

Overview of HPD

The HPD (HumanDetector) is a **tile based convolutional network** designed to detect humans in images. Its key characteristics:

1. **Inputs:** $x \in \mathbb{R}^{N \times 1 \times H \times W}$ grayscale images.
2. **Tiling:** Divide image into $\text{tile_rows} \times \text{tile_cols}$ patches without overlaps.
3. **Per tile Feature Extraction:** The use of LBCLayer to perform sparse and task-aware feature extraction from each patch.
4. **Features Aggregation / Stitching:** Features from all the patches aggregated into one feature vector.
5. **Classifier Head:** The output dense layers that provide logit values for num_classes . (E.g., Human vs. Not Human).

Mathematically:

$$f_{\theta}(x) = \text{Classifier}(\text{Concat}_{i,j}(\text{TileExtractor}(x_{i,j}))) \quad (4-5)$$

Where $x_{i,j}$ is the (i, j) -th tile.

Tiling Stage

The tiling stage prepares the input for localized processing.

- **Input:** $x \in \mathbb{R}^{N \times 1 \times H \times W}$
- **Parameters:** $\text{tile_rows}, \text{tile_cols}$
- **Tile size:**

$$h = \frac{H}{\text{tile_rows}}, \quad w = \frac{W}{\text{tile_cols}} \quad (4-6)$$

- **Split input into tiles:**

$$x_{i,j} = x[:, :, i \cdot h : (i + 1) \cdot h, j \cdot w : (j + 1) \cdot w] \quad (4-7)$$

Purpose:

- Localizes feature extraction $\rightarrow$ each tile acts as a **node** (MCU-like processing).
- Improves **spatial granularity** for detecting humans.
- Tiles are **independent**, can be processed in parallel.

Tile Feature Extractor (LBC-based)

Every tile goes through *tile_extractor* whose structure is described below:

- **LBC Layer 1:** 1to16 channels with sparsity = 0.9.
- **MaxPool2d:** Kernel size 2 by reducing spatial dimensions by half.
- **LBC Layer 2:** 16to32 channels with sparsity = 0.9.
- **AdaptiveAvgPool2d:** pooling spatial dimensions to 1×1 .

For individual tile $x_{i,j}$:

$$x' = \text{LBCLayer}(x_{i,j}) = \text{Conv1x1}(\text{ReLU}(\text{Conv}_{\text{binary}}(x_{i,j}))) \quad (4-8)$$

Aggregation / Stitching

After feature extraction, the tile outputs get aggregated together:

1. **Flatten the output for each tile:** $f_{i,j} \in \mathbb{R}^{32}$
2. **Concatenate all tiles:** If number of tiles = $\text{tile_rows} \times \text{tile_cols}$:

$$f_{\text{combined}} = [f_{1,1}, f_{1,2}, \dots, f_{\text{tile_rows}, \text{tile_cols}}] \in \mathbb{R}^{32 \cdot \text{num_tiles}} \quad (4-9)$$

This process can be considered analogous to graph aggregation, treating every tile as a graph node

Classifier Head

The classifier performs an MLP on the concatenated vector:

- **Input:** $f_{\text{combined}} \in \mathbb{R}^{32 \cdot \text{number of tiles}}$
- **Layers:** $\text{Linear}(32 \cdot \text{number of tiles} \rightarrow 64) \rightarrow \text{ReLU} \rightarrow \text{Linear}(64 \rightarrow \text{num_classes})$
- **Output Classification Logits:**

$$\hat{y} \in \mathbb{R}^{N \times \text{num_classes}} \quad (4-10)$$

Uses **CrossEntropyLoss** when training:

$$L = - \sum_{c=1}^{\text{num_classes}} y_c \log \text{softmax}(\hat{y}_c) \quad (4-11)$$

Computational Efficiency

The LBCLayer reduces Floating Point Operations (FLOPs) significantly:

$$\text{FLOPs} \approx C_{in} \cdot C_{out} \cdot k^2 \cdot H \cdot W \quad (4-12)$$

- **Binary conv:** Multiplications are basically simple additions/subtractions.
- **1x1 conv:** Provides lightweight mixing of channels.
- **Tile based processing:** Can be parallelized across multiple nodes.

Mathematical Summary of Forward Pass

1. Tile splitting:

$$\{x_{i,j}\}_{i,j=1}^{R,C} = \text{Tile}(x) \quad (4-13)$$

2. Feature extraction (per tile):

$$f_{i,j} = \text{AdaptiveAvgPool2d}(\text{LBCLayer2}(\text{MaxPool2d}(\text{LBCLayer1}(x_{i,j})))) \quad (4-14)$$

3. Flatten and concatenate:

$$f_{\text{combined}} = \text{Concat}([f_{i,j}]) \quad (4-15)$$

4. Classification:

$$\hat{y} = \text{Linear}_2(\text{ReLU}(\text{Linear}_1(f_{\text{combined}}))) \quad (4-16)$$

5. Loss:

$$L = - \sum_{c=1}^{\text{num_classes}} y_c \log \text{softmax}(\hat{y}_c) \quad (4-17)$$

Advantages of HPD Architecture

1. **Sparse and efficient:** LBC reduces compute and memory footprint.
2. **Tile based learning:** Learns local patterns for small human parts.
3. **Flexible:** Can adapt tile size and sparsity based on resources.
4. **Parallelization:** Processed independently per tile.
5. **Low data scenario:** Set anchors minimize training needs.

4.2. Hardware Requirements and system configuration

The design employs a distributed architecture that is multi-node based on the ESP32 and ESP32-S3 MCU. In order to maximize efficiency and take advantage of limited resources, the system is structured hierarchically in three levels.




Figure 4-3. System Architecture showing main components and their interactions.




Figure 4-4. Detailed Data Flow Diagram of the Tiling and Inference Pipeline.

4.2.1. Input Layer

Centered around an ESP32 CAM [?], this layer manages high bandwidth data acquisition and initial pre processing. Its primary role is to prepare and distribute workloads to the processing tier through the following pipeline:

Luminance Conversion (Grayscale)

The processing of three-channel data requires higher memory buffers ($O(3n)$ vs $O(n)$). This frame will have to be translated into an 8-bit grayscale format using the luminance weight formula as specified in equation (4-18) [?].

$$P(y = 1 | x) = \frac{e^{z_{\text{person}}}}{e^{z_{\text{bg}}} + e^{z_{\text{person}}}} \quad (4-18)$$

Where z_{person} and z_{bg} represents the logits corresponding to *person* and *background* class, respectively. Data compression reduces the input data volume by 66.6%, creating more room for heap memory allocation for tiling and transmission buffers.

$$Y = 0.299R + 0.587G + 0.114B \quad (4-19)$$

Partitioning Distribution

Tiling stage involves the use of the Tiling module to select n overlapping tiles from the input image. The importance of this stage cannot be underestimated since it is vital in ensuring that important high-resolution information is retained within the resource-limited computing system. With an overlap ratio of 25%, this helps preserve humans located on the boundary of two tiles.

4.2.2. Communication Protocol Implementation

The Espressif Systems (ESP) protocol will be the communication backbone [?], Selected due to high efficiency in transferring data through peer-to-peer interaction and low latency feedback. Though ESP is efficient, its bandwidth is limited up to a maximum of 771 Kilobyte (KB)/s. In this regard, efficient packing of data becomes crucial since the throughput capacity will dictate how fast the tile data transfer takes place while the results packets will comply with the packet size constraints less than 250 bytes. It is important that data transfer occurs quickly since the gains from parallel processing could be offset by the time required for communicating.

4.2.3. Feature Extraction Layer

- **Parallel Processing:** Each node receives a specific low resolution tile from the Input Layer.
- **Local Inference:** Each node runs the HPD model on its assigned tile, extracting features and computing an attention weight that indicates the likelihood of human presence in that tile.
- **Output Generation:** Each node outputs a feature map and an attention weight, which are then transmitted to the Output Layer for aggregation and final decision-making.

4.2.4. Output Layer (Aggregation & Logic)

This layer is the last stage of the architecture and it is the core decision making layer. The feature map and the attention weight of each tile are taken from the feature extractor nodes, and the following operations are performed:

- **Attention score calculation:** The attention score for each tile is computed using the SoftMax function, which normalizes the quantized scores into a probability distribution. This allows the system to make weighted final decision.
- **Weighted feature aggregation:** Using the attention scores, we do a weighted sum of the feature maps from all tiles. This effectively combines the information, while giving more weight to tiles that are more likely to contain human features.
- **Final Classification:** The aggregated feature map is then passed to a classifier layer (e.g., fully connected layer) to compute a confidence score for the presence of a human in the whole scene. This score is then used to make a binary decision (human detected or not) that can be displayed on an output device such as an LCD screen.

4.3. Hardware Used

The core hardware of this project is based on the MCUs of the ESP32-family. It is chosen because they are low cost, low power, and can run edge-level machine learning tasks. This system uses a distributed architecture. Here, image capturing, processing and decision-making are separated across different ESP32 devices to overcome memory and processing limitations of a single microcontroller. This hardware design allows the system to perform real-time HPD without relying on cloud servers, making it faster, more private, and more energy efficient.




Figure 4-5. Wiring Interface between Programming (FTDI), Processing (ESP32-S3), and Visualization (I2C LCD) Units.

The hardware architecture as shown in 4-5 consists of a distributed processing network designed for real-time data acquisition and visualization. The system is partitioned into three primary functional blocks: the Sensing Node (ESP32-CAM), the Processing Cluster (ESP32-S3 units), and the Output Interface (I2C LCD connected via ESP32)

4.3.1. ESP32-CAM

The ESP-CAM module is used as the input hardware of the system. Its main role is to capture visual data from the environment. Here, the camera continuously captures images of 224×224 at a controlled rate of 1 every 15 seconds. Since color images require more memory and computation, the ESP-CAM converts the captured rgb images into grayscale format. After this, the high-resolution image is divided into smaller tiles of 96×96 with overlap of 25% between adjacent tiles. These smaller image tiles are easier to process and then transmit to other ESP devices using ESP-NOW protocol for further processing. Here, the tiles are evenly distributed between the two esp32s3 devices when the total number of tiles is even; if the number of tiles is odd, the first node is assigned one additional tile.

4.3.2. USB to TTL converter (FTDI)

The USB to TTL (FTDI) converter is used to connect the ESP32-CAM boards to a laptop for programming. It is because ESP32-CAM does not have a built-in USB interface for programming or serial communication. The FTDI module acts as a bridge between them since a laptop communicates using USB and the ESP32 uses TTL serial (UART) signals. In this project, the FTDI also supplies 5V power from the laptop via USB, which removes the need for an external power source.

Connect GPIO0 to GND in ESP32-CAM for flashing. This indicates that the ESP32-CAM is ready to receive new code. After uploading the code, we can disconnect GPIO0 from GND.

Table 4-1. Wiring the FTDI to ESP32-CAM

FTDI	ESP32-CAM
VCC	5V
GND	GND
RX	U0T
TX	U0R

4.3.3. ESP32-S3

The ESP32-S3 MCUs act as the processing or inference nodes in the system. Each ESP32-S3 receives a small image tile from the ESP32-CAM, on which it runs a Tiny Machine Learning (TinyML)-based HPD model. These boards are selected because they are optimized for AI workloads and support low-bit operations. These are necessary to run neural networks with limited memory. Each ESP32-S3 processes its assigned tile independently and decides if human features are present. Each ESP32-S3 receives a small image tile from the ESP32-CAM, on which it runs a Tiny Machine Learning

(TinyML)-based HPD model. These boards are selected because they are optimized for AI workloads and support low-bit operations. These are necessary to run neural networks with limited memory. Each ESP32-S3 processes its assigned tile independently and decides if human features are present. Instead of sending large image data back, it sends only a compact result, such as a confidence score and tile location, which minimizes communication overhead.

4.3.4. ESP32

The ESP32 board is used as the aggregation node of the system. It is used to collect the results from all ESP32-S3 and uses attention weight from all tiles to compute attention score using SoftMax function and then uses the attention score to perform weighted sum of feature maps of tiles. Then, it passes the new feature map through classifier layer to compute the confidence score. This confidence score is to form a final decision about human presence in the entire scene.

4.3.5. I2C LCD Display

In this system, the I2C LCD display is used to show the final output of the project. After the ESP32-S3 boards finish processing the image tiles and send their results, the ESP32 combines all the information and decides whether a human is present or not. This final decision is then shown on the LCD screen as **“Human Detected”** or **“No Human Detected”**. The I2C LCD has been chosen because it needs only two wires for communication. It keeps the wiring simple and saves GPIO pins on the ESP32.

Table 4-2. Wiring the I2C LCD to ESP32

I2C LCD Pin	ESP32 Pin	Purpose
VCC	3.3V or 5V	Power Supply
GND	GND	Ground Connection
SDA	GPIO21	Data Line
SCL	GPIO22	Clock Line

4.3.6. Power Supply

All hardware components in the system are powered through USB connections using a laptop. During development and testing, the laptop provides stable 5V power supply to the ESP32-CAM, ESP32-S3 and ESP32 boards.

4.3.7. Low Bit Quantization Implementation

The baseline operational model will use Post training Quantization (PTQ) to convert the HPD model to symmetric 8 bit Integer (INT8) weights and asymmetric INT8 activations. Further optimization will involve the experimental deployment of 4-bit Integer (INT4)

quantization for specific non critical layers. This particular strategy will bring the use of memory to its barest minimum and achieve an ultimate Flash utilization that is less than 1.0 Megabyte (MB). This will require us to compensate for precision degradation through use of calibration data as well as equalization of the weights within each layer to minimize precision degradation to a level that is not more than 90% of the pruned 32-bit Floating Point 32-bit Floating Point (FP32). The justification for this strategy is that it allows creation of a reliable model and fits within 512 KB SRAM.

4.4. Evaluation Metrics and Performance Framework

4.4.1. Detection Reliability Matrix (The Deep Learning Layer)

The main goal is to fix the stationary detection issue that PIR sensors have. The following standard computer vision criteria will be used to assess our CNN model:

Recall (Sensitivity)

Measured especially for participants who are motionless to demonstrate the system's capacity to sustain presence signals in situations where movement is minimal. The Equation is as shown in equation (3-3):

$$Recall = TruePositive(TP)/(TP + FalseNegative(FN)) \quad (4-20)$$

Precision

To ensure environmental noise (shadows, HVAC airflow) does not trigger false positives in the model's inference. The Equation is as shown in equation (3-4):

$$Precision = TP/(TP + FalsePositive(FP)) \quad (4-21)$$

4.4.2. Distributed System Efficiency Matrix (The Computing Layer)

We must measure the benefits of distributing the inference burden among several ESP32-S3 nodes using the Distributed Adaptive Tiling (DAT) framework as this project makes use of a Distributed Cluster.

End-to-End Latency (total): The duration between the Master node's picture acquisition and the ultimate detection outcome.

Memory Footprint: Maximum SRAM used by any single node during inference.


Communication Delay: The proportion of the entire cycle used for local Deep learning computation as opposed to ESP tile transmission.

5. RESULTS AND ANALYSIS

The following report details the training and evaluation of a LBC architecture for HPD. The model was trained over 45 epochs using a *OneCycleLR* scheduler before early stopping was triggered.

5.1. MobileNetV2 Architecture Performance

The utilization of a MobileNetV2 architecture for the task of HPD demonstrated significantly higher accuracy compared to the lightweight LBCNN baseline, at the cost of increased model complexity. The model was trained for 12 epochs before early stopping was triggered, using a *OneCycleLR* scheduler with a maximum learning rate of 0.01.



(a) Training and validation loss over 45 epochs. (b) Normalised confusion matrix on the test set.

Figure 5-1. MobileNetV2 training dynamics and test-set evaluation results.

The model was trained on a significantly larger dataset of 20,736 samples (split into 85% training, 5% validation, and 10% test) compared to the LBCNN baseline. As shown in Figs. 5-1a and 5-1b, both training and validation loss decrease consistently across epochs with strong convergence, indicating that the larger dataset effectively mitigated the overfitting observed in the LBCNN model. The best validation accuracy of **87.34%** was recorded at epoch 7, and early stopping was triggered at epoch 12 with a patience of 5.

Table 5-1. MobileNetV2 Training Run Summary

Metric	Value	Epoch
Best validation accuracy	87.34%	7
Best validation loss	0.3145	6
Final training accuracy	89.32%	12
Test accuracy	86.18%	—
Test Area Under the Curve (AUC) (ROC)	0.931	—
Early stopping	—	12

5.2. MobileNetV2 Test-Set Evaluation


Following training, the best-checkpoint model was evaluated on a held-out test set of 2,432 images. Table 5-2 presents the per-class precision, recall, and F1-score.

Table 5-2. Per-Class Classification Report on the Test Set ($n = 2432$) for MobileNetV2


Class	Precision	Recall	F1-score	Support
Class 0 (No human)	0.7958	0.8955	0.8427	1005
Class 1 (Human)	0.9193	0.8381	0.8768	1427
Macro average	0.8575	0.8668	0.8598	2432
Weighted average	0.8682	0.8618	0.8627	2432
Overall test accuracy				86.18%

5.3. LBCNN Architecture Performance

The utilization of LBC architecture for the task of human presence detection demonstrated the feasibility of a lightweight model with reliable accuracy on resource constrained MCUs.



(a) Training and validation loss over 45 epochs.



(b) Normalised confusion matrix on the test set.

Figure 5-2. LBCNN training dynamics and test-set evaluation results.

As illustrated in Figs. 5-2a and 5-2b, the LBCNN model shows consistent reduction in training loss over 45 epochs, decreasing steadily from 0.70 to approximately 0.51, while the validation loss remains highly volatile throughout training with sharp spikes reaching as high as 1.10, a behaviour attributed to the limited dataset size of 737 training samples where individual batches exert disproportionate influence on the validation metric. The normalised confusion matrix indicates broadly balanced per class recall of 0.72 for Class 0 and 0.71 for Class 1, indicating that the weighted sampling strategy was successful in avoiding majority class collapse. The 28% false alarm rate and 29% missed detection rate suggest that the binary convolutional features lack sufficient discriminative precision to separate the two classes cleanly. It is a direct consequence of the constrained model capacity and limited training data.

Table 5-3. Training Run Summary

Metric	Value	Epoch
Best validation loss	0.5721	30
Best validation accuracy	77.17%	43
Final training accuracy	78.56%	45
Test accuracy	71.74%	—
Test AUC (roc)	0.777	—
Early stopping	—	45

5.4. Test-Set Evaluation

Following training, the best-checkpoint model was evaluated on a held-out test set of 92 images. Table 5-4 presents the per-class precision, recall, and F1-score.

Table 5-4. Per-Class Classification Report on the Test Set ($n = 92$)

Class	Precision	Recall	F1-score	Support
Class 0 (No human)	0.6190	0.7222	0.6667	36
Class 1 (Human)	0.8000	0.7143	0.7547	56
Macro average	0.7095	0.7183	0.7107	92
Weighted average	0.7292	0.7174	0.7203	92
Overall test accuracy				71.74%

5.5. Comparative Analysis: LBCNN vs. MobileNetV2

Table 5-5 summarises the key performance and resource metrics for both architectures studied in this work. The LBCNN model is intended for direct deployment on resource-constrained MCUs while MobileNetV2 is a higher capacity reference baseline trained on a larger dataset.

Table 5-5. Performance Comparison: LBCNN vs. MobileNetV2

Metric	LBCNN	MobileNetV2
Best validation accuracy	77.17%	87.34%
Best validation loss	0.5721	0.3145
Final training accuracy	78.56%	89.32%
Test accuracy	71.74%	86.18%
Test AUC (ROC)	0.777	0.931
Class 0 Precision	0.619	0.796
Class 0 Recall	0.722	0.896
Class 0 F1-score	0.667	0.843
Class 1 Precision	0.8000	0.9193
Class 1 Recall	0.7143	0.8381
Class 1 F1-score	0.7547	0.8768
Number of parameters	19,779	2,225,858
Model size (FP32)	80 KB	8490 KB
8-bit Integer (INT8) Model size	20 KB	2122 KB

5.6. Overall Comparative Analysis

The experimental results summarised in Table 5-5 show that the MobileNetV2 outperforms the LBCNN in all the evaluation metrics with a clear edge in performance. The MobileNetV2 achieved 86.18% test accuracy and 0.931% glsauc, while the LBCNN achieved 71.74% and 0.777%, respectively, which represents an improvement of 14.44% percentage points and 0.154%. This gap is present for all the per-class metrics, with MobileNetV2 being better than glslbcnn in terms of both precision and recall for both classes.

The difference in training behaviour is equally clear. The glslbcnn loss and accuracy curves are very volatile for all the 45 epochs with the validation accuracy jumping sharply between successive epochs. The instability is a direct result of the very small training set (737 samples) where individual batches have an outsized impact on the gradient signal. The MobileNetV2 model trained on 18,432 samples with ImageNet pretrained weights converges smoothly in 12 epochs, with the loss and accuracy curves declining and increasing monotonically, respectively. The pretrained initialisation removes the need to learn low level features from scratch, and the larger dataset provides stable gradients together explaining why MobileNetV2 converges faster, more cleanly and to a significantly better optimum..

This picture is supported by the roc and precision–recall curves. The roc curve of MobileNetV2 goes up sharply towards the top left corner, suggesting high true positive rates for low false positive rates and demonstrating that the class probability scores are well separated. The LBCNN ROC curve is more gradual, reflecting that its 64-dimensional binary feature representations achieve less class separability. Similarly, MobileNetV2 achieves higher precision over a wider range of recall thresholds. LBCNN precision drops more quickly with increasing recall, in line with a model working close to the limits of its representational capacity.

However, this comparison should be understood in the context of the fundamentally different constraints under which each model operates. MobileNetV2 has 2,225,858 parameters and takes up 8.49 MB in FP32, which is way beyond the memory budget of any commodity MCU. The LBCNN, which has only 19,779 parameters and a deployed size of just 0.02 MB after INT8 quantization, easily fits within the ESP32-S3’s 512 KB SRAM. The 14.44 percentage point accuracy trade-off is thus not a shortcoming but an intentional result of designing for deployment on hardware that is about **425× more memory-constrained** than what MobileNetV2 requires. For real-time, battery-powered, GPU-free edge inference, LBCNN offers a principled and practical solution to the accuracy-efficiency trade-off of Tiny Machine Learning (TinyML) deployments.

5.7. System Overview and Evaluation Methodology

We evaluate this distributed inference pipeline for four different processing stages: frame capture and tiling at the ESP-CAM, tile transmission from the ESP-CAM to the two ESP32-S3 worker nodes, tile-level feature extraction and attention scoring at each worker node, weighted feature transmission from the worker nodes to the ESP32 aggregator, and final classification at the aggregator. Each stage was independently benchmarked against dedicated test sketches with synthetic inputs to isolate stage specific latency from end-to-end confounding factors. All the timing measurements were done using the ESP32 hardware *micros()* counter, with the CPU clock set to 240MHz. Table 5-6 presents the complete latency breakdown across all pipeline stages.

Table 5-6. End-to-End Pipeline Latency Breakdown

Stage	Node	Detail	Latency
Frame capture + tiling	ESP-CAM	224×224 → 9 tiles	4 ms
Tile transmission (CAM → Workers)	Worker 1	5 tiles × 38 fragments	3,243 ms
Tile transmission (CAM → Workers)	Worker 2	4 tiles × 38 fragments	1,183 ms
Tile inference (extractor + attention)	Worker 1	5 tiles sequential	13,284 ms
Tile inference (extractor + attention)	Worker 2	4 tiles sequential	10,627 ms
Feature transmission (Workers → Aggregator)	Per tile avg	69-byte payload	472 ms
Aggregator (softmax + sum + classifier)	ESP32	9 tiles aggregated	0.342 ms
Total pipeline (critical path)	Worker 1 bottleneck		~16,999 ms

5.7.1. ESP-CAM: Frame Capture and Tiling

The ESP-CAM node captures a 224×224 gray-scale frame and splits it into nine overlapping 96×96 tiles with a step size of 72 pixels (25% overlap) along both spatial axes. The whole capture-to-tile pipeline takes 4 ms, rendering this stage completely insignificant compared to all downstream processing. This confirms the fact that the tiling strategy itself incurs no meaningful overhead, and that the ESP-CAM is able to achieve the target 1 FPS frame rate from a capture perspective alone.

The 4 ms figure encompasses camera sensor readout, RGB565-to-grayscale conversion, and the nine sequential tile extraction operations. At this rate, the ESP-CAM could

theoretically support up to 250 FPS in isolation, meaning the capture stage will never be the bottleneck regardless of how the downstream pipeline is optimised.

5.7.2. Tile Transmission: ESP-CAM to Worker Nodes

The first latency stage in the pipeline is the transmission of tiles from the ESP-CAM to the two worker nodes via ESP-NOW. Each tile is 96×96 , and is 9,216 bytes in size. ESP-NOW has a maximum frame size of 250 bytes, so each tile is fragmented into 38 sequential packets of 247 data bytes each, with a synchronous ACK-wait between fragments for flow control.

Worker 1 was assigned tiles 0—4 (five tiles) and took 3,243 ms to transmit. Worker 2 was assigned tiles 5—8 (four tiles) and took 1,183 ms. The difference between these two numbers 3,243 ms vs. 1,183 ms for one more tile is disproportionate and deserves analysis

The main reason for this is channel contention. Both workers are on the same 2.4 GHz ESP-NOW channel and the ESP-CAM first transmits to Worker 1 then switches to Worker 2. Worker 1 then takes the full serialised cost of $5 \times 38 = 190$ fragments with ACK-wait on each, but Worker 2 gets a cleaner channel once Worker 1's transmission is over. Worker 1 has a about 648 ms per tile transmission cost ($3,243 / 5$). Worker 2 has a about 296 ms per tile transmission cost ($1,183 / 4$). The reason for this 2.2 $\times$ difference is the compounding effect of the sequential ACK-wait under a loaded channel, and not an inherent difference in channel quality between the two nodes.

Both of these figures are well over the 1,000 ms frame budget target. The transmission stage alone is 4.4 $\times$ over budget, as the ESP-CAM takes about 4,426 ms ($3,243 + 1,183$) total time to finish sending all nine tiles. This indicates that the most promising area of optimisation after inference latency is tile fragmentation and sequential ACK-based flow control.

5.7.3. Worker Node Inference: Tile Extractor and Attention

Each ESP32-S3 worker node runs two TFLite Micro models sequentially per tile: the tile extractor (60 KB INT8, producing a 64-dimensional feature vector) and the attention head (6 KB INT8, producing a scalar logit). Both models execute from a 600 KB arena allocated in Octal PSRAM via *heap_caps_aligned_alloc* with 16-byte alignment.

Worker 1 processed five tiles sequentially, completing in 13,284 ms total, yielding an average per tile inference time of 2,657 ms. Worker 2 processed four tiles sequentially, completing in 10,627 ms total, yielding an average per tile inference

time of 2,657 ms consistent with Worker 1, confirming that per tile inference time is stable and node-independent.

The major source of such latency is the performance of PSRAM read operations. The required size of the arena is 451 KB which surpasses the internal SRAM capacity of the ESP32-S3 board available after system allocations (~180 KB). Consequently, all intermediate activations tensors have to be stored in Octal PSRAM which introduces higher latency compared to accessing internal SRAM because of the higher latencies of accessing Octal PSRAM via the SPI bus and the penalty for missing the cache. Accessing SRAM on the ESP32-S3 runs at the same rate as the CPU (~4 ns on 240 MHz), while PSRAM access introduces a latency of about 20 to 40 clock cycles resulting in the observed ~2.66 s per-tile compared to 100 ms when using SRAM. Table 5-7 presents the per tile inference breakdown.

Table 5-7. Worker Node Per-Tile Inference Latency

Metric	Worker 1 (5 tiles)	Worker 2 (4 tiles)
Total inference time	13,284 ms	10,627 ms
Avg per-tile inference	2,657 ms	2,657 ms
Tile extractor arena used	441 KB (PSRAM)	441 KB (PSRAM)
Attention arena used	<1 KB (SRAM)	<1 KB (SRAM)
Model size (extractor)	60 KB	60 KB
Model size (attention)	6 KB	6 KB
Quantisation	INT8 (fully quantised)	INT8 (fully quantised)
Float32 tensors remaining	0	0

5.7.4. Feature Transmission: Worker Nodes to Aggregator

Each worker node transmits a 69 byte *WorkerPayload* struct to the aggregator upon completing inference for each tile. The payload contains the tile ID (1 byte), raw attention logit (4 bytes as float32), and the quantised 64 dimensional feature vector (64 bytes as int8). At 69 bytes, this fits comfortably within a single ESP NOW frame (250 byte maximum), requiring no fragmentation.

The average one way transmission delay per payload was measured at 472 ms using a round trip ACK scheme with one way delay estimated as $RTT/2$. This figure is substantially higher than the theoretical ESP NOW unicast latency of approximately 2 to 5 ms and warrants explanation.

The 472 ms figure reflects the queued transmission delay rather than raw channel latency. Because each worker node processes tiles sequentially and inference takes ~2,657 ms

per tile, payloads accumulate in the transmission queue during inference. The measured 472 ms average therefore includes both the actual ESP NOW channel time (~ 2 to 5 ms) and the queuing delay imposed by the serialised inference pipeline. The true channel transmission time for a 69 byte ESP NOW packet is sub millisecond; the observed 472 ms is a system level metric reflecting the pacing of the inference pipeline rather than a channel limitation.

Total transmission time for all nine feature payloads (five from Worker 1 plus four from Worker 2) is approximately 4,248 ms (9×472 ms), though in practice Worker 1 and Worker 2 transmit in parallel once their respective inference completes, reducing the effective wall clock transmission time to approximately $\max(5, 4) \times 472$ ms = 2,360 ms under ideal parallel operation.

5.7.5. Aggregator Node: Softmax, Weighted Sum, and Classification

The ESP32 aggregator node receives nine *WorkerPayload* packets, performs softmax normalisation across the nine raw attention logits, computes the attention-weighted sum of the nine 64-dimensional feature vectors, and passes the resulting aggregated feature through the INT8 classifier to produce a binary label.

The complete aggregator pipeline executed in 0.342 ms, making it the fastest stage in the entire system by four orders of magnitude. Table 5-8 presents the aggregator latency breakdown.

Table 5-8. Aggregator Node Processing Latency

Sub-stage	Latency
Softmax (9 logits) + weighted sum (9×64)	~ 0.330 ms
Classifier inference (INT8, $64 \rightarrow 32 \rightarrow 2$)	~ 0.012 ms
Total aggregator delay	0.342 ms
Classifier model size	4 KB
Classifier arena	<32 KB (SRAM)

The negligible aggregator latency confirms that the architectural decision to centralise only the lightweight aggregation and classification operations at the aggregator node rather than any feature extraction is well-founded. The 4 KB INT8 classifier runs entirely from internal SRAM with no PSRAM dependency, and the softmax and weighted sum operations involve only 9 scalar values and $9 \times 64 = 576$ multiply-accumulate operations, well within the ESP32's capability at microsecond timescales.

5.7.6. End-to-End Pipeline Analysis

Table 5-9 presents the complete pipeline latency on the critical path (Worker 1 as bottleneck) alongside the 1,000 ms frame budget target.

Table 5-9. Critical Path Latency vs. Target Frame Budget

Stage	Latency	% of Budget
Frame capture + tiling	4 ms	0.4%
CAM → Worker 1 transmission	3,243 ms	324.3%
Worker 1 inference (5 tiles)	13,284 ms	1,328.4%
Worker 1 → Aggregator transmission	2,360 ms	236.0%
Aggregator processing	0.342 ms	0.03%
Total critical path	~18,891 ms	1,889%
Target frame budget	1,000 ms	100%

The end-to-end pipeline currently operates at approximately 18.9× over the 1,000 ms frame budget, with Worker 1 inference (13,284 ms) being the dominant bottleneck, contributing 70.3% of total pipeline latency. Tile transmission from the ESP-CAM to the workers (3,243 ms) contributes a further 17.2%, and feature transmission from workers to the aggregator contributes 12.5%. The capture and aggregation stages together account for less than 0.03% of total latency and are effectively free in comparison.

5.7.7. Quantisation and Model Size Analysis

Table 5-10 summarises the quantisation outcomes for all three deployed models.

Table 5-10. Model Quantisation Summary

Model	FP32 Size	INT8 Size	Reduction	Float32 Tensors	Arena Required
Tile extractor	~155 KB (ONNX)	60 KB	2.6×	0	451 KB (PSRAM)
Attention head	~11 KB (ONNX)	6 KB	1.8×	0	<1 KB (SRAM)
Classifier	~9 KB (ONNX)	4 KB	2.3×	0	<32 KB (SRAM)

All three models were verified to contain zero float32 activation tensors following the two-stage conversion pipeline (ONNX → TF SavedModel → TFLite INT8 via representative dataset calibration). The absence of hybrid quantisation where some

layers remain in float32 was confirmed by iterating over all tensor details in the TFLite interpreter and checking quantisation scale parameters, resolving an earlier issue where *onnx2tf*'s *-oiqt* flag produced hybrid models that caused *Hybrid models are not supported* errors on TFLite Micro.

The tile extractor's arena requirement of 451 KB despite the model itself being only 60 KB reflects TFLite Micro's greedy memory planning strategy, which allocates scratch buffers for intermediate activations across the three *LBCLayer* blocks. Each block processes feature maps of spatial dimensions up to 96×96, requiring substantial intermediate storage that cannot be released until the layer completes.

5.7.8. Discussion: Performance Gap and Path to Real-Time Operation

The results show a clear and quantifiable performance gap between the current prototype and the real time target. The system is architecturally correct: all pipeline stages work as intended, data flows correctly between nodes, INT8 inference produces valid outputs, and the aggregator correctly fuses nine partial representations into a binary decision. However, the system currently does not meet the latency requirements to operate in real-time at 1 FPS. The results clearly point the way to close this gap. To reduce per tile inference time from 2657 ms to <150 ms, either the tile extractor arena needs to fit entirely within internal SRAM (a model with arena requirement of <~ 180 KB can be achieved by reducing the channel count or by simplifying the architecture), or a hardware platform with faster external memory such as the ESP32-P4 (with its on chip cache backed PSRAM) needs to be used. Meanwhile, the 3,243 ms Worker 1 transmission bottleneck could be addressed by swapping out the sequential ACK based fragmentation scheme for a broadcast based transmission or by increasing the ESP-NOW payload limit via custom fragmentation. The aggregator node, with its 0.342 ms processing time and 4 KB model footprint, confirms that the fusion and classification stage of the architecture imposes essentially zero overhead and would remain viable even as the upstream pipeline is accelerated by one to two orders of magnitude.

6. CONCLUSION

6.1. Achievements

This project addressed a fundamental constraint in edge deployed computer vision: the inability of resource limited microcontrollers to process high resolution imagery without destroying the spatial detail necessary for reliable detection. The proposed solution, a distributed tile based inference pipeline for HPD demonstrated that spatially decomposed neural inference across a heterogeneous network of commodity ESP32 class hardware is not only architecturally feasible but implementable end to end without GPU acceleration, cloud offloading, or specialised neural processing hardware.

The system successfully achieved the following concrete outcomes. A complete model training and quantisation pipeline was developed, converting a custom Improved LBCNN from floating-point PyTorch to fully INT8 TensorFlow Lite Micro format, achieving approximately 4× model size reduction while preserving functional inference capability on the target hardware. The tile extraction, attention based feature aggregation, and binary classification pipeline was validated end to end in simulation, with per tile inference, transmission delay, and aggregator latency each measured independently through dedicated benchmark sketches. The distributed communication layer was implemented using the ESP-NOW protocol, enabling sub millisecond peer to peer transmission of 69 byte weighted feature payloads between worker nodes and the aggregator. Fragment based tile reassembly with duplicate suppression was implemented and validated on the ESP32-S3 worker nodes, demonstrating robustness to packet reordering under the 250 byte ESP-NOW frame size constraint. The full quantisation pipeline from PyTorch checkpoint through ONNX, TF SavedModel, and TFLite INT8 conversion was automated into a single reproducible script, with all three sub module headers (tile extractor, attention, classifier) verified to contain zero float32 tensors prior to deployment.

6.2. Limitations

Several constraints were encountered during implementation that bound the current system's performance and generalisability.

Inference latency. The dominant bottleneck in the system is the per tile inference time on the ESP32-S3 worker nodes, measured at approximately 2.68 seconds per tile when the TFLite arena is allocated in Octal PSRAM. This far exceeds the 1,000 ms target frame period, meaning the current implementation cannot sustain real time throughput at the intended 1 FPS capture rate. The root cause is the high latency of PSRAM access relative to internal SRAM: the tile extractor requires approximately 451 KB of arena

memory exceeding the ESP32-S3's available internal SRAM forcing all intermediate activations through the slower PSRAM bus. This is an architectural constraint of the chosen model and hardware compatibility rather than a software defect.

Model accuracy. The reported validation accuracy of 35.87% for the LBCNN backbone is not good enough for deployment in a real detection system. This suggests the use of synthetic calibration data during INT8 quantisation instead of real validation imagery and a training dataset that might not be sufficiently representative of the diversity of human subjects and environmental conditions encountered in practice. The quantisation pipeline and deployment infrastructure are solid. The accuracy limitation is a function of the training data and training duration, not the architectural approach.

Hardware availability. The absence of an ESP32-CAM during development prevented the full end to end pipeline from live frame capture, tiling, distributed inference and aggregated classification to be validated on real imagery. All tile inputs to the worker nodes during benchmarking were synthetic, generated programmatically to exercise the inference and communication pathways. System behavior on real world frames with real spatial variation has not been validated.

Fragment reassembly reliability. The current tile fragmentation scheme sends 38 sequential ESP-NOW packets per tile with a simple ACK wait flow control. Under network congestion or when both workers transmit simultaneously, packet collisions at the aggregator's receive buffer can cause fragment loss and incomplete tile reassembly. While duplicate suppression is implemented, a formal retransmission mechanism is absent.

Single channel input. The model operates on grayscale (single channel) input, which discards chrominance information that could assist in distinguishing human subjects from similarly shaped environmental objects in colour imagery.

6.3. Future Work

Several concrete directions emerge from this work that would substantially improve the system's performance, robustness, and real world applicability.

Reducing inference latency. The most important step needed for the system to work in real time would be to lower the per tile inference time to under 200 ms. Three different strategies are worthy of investigation. Firstly, compression of the LBCNN backbone using structured pruning would lower the amount of channels in each layer, thus lowering the arena size to fit into internal SRAM and removing the need to pay PSRAM latencies. Secondly, changing the LBCNN backbone for a version of MobileNetV3(Small) or EfficientNet(Lite) optimised for TFLite Micro could provide similar feature detection

capabilities while using much less memory. Lastly, the dual core structure of the ESP32-S3 could be used for running fragment reassembly and inference in a pipeline fashion on cores 0 and 1 respectively. **Improving model accuracy.** A validation accuracy of 35.87% is achieved during the proof of concept training session. A systematic retraining process with an organised set of data related to human detections from datasets like INRIA Person Dataset or Crowd Human, along with appropriate augmentation, class balancing, and additional training, will be the quickest route towards achieving an accuracy score that can be implemented into a model. Moreover, using actual images instead of randomly generated Gaussian noise for INT8 quantisation will lead to better results during quantisation and less accuracy drop from Float to INT8. **Real ESP-CAM integration.** It is most critical to validate the entire pipeline using a live ESP32 CAM stream. Doing so will provide insights into any timing issues that arise from interactions between capture, fragmentation, transmission, and inference, which cannot be identified within a synthetic benchmark test environment. Furthermore, it will indicate whether the approach of 25% overlap works well for tile-boundary subjects.

Reliable fragment transport. A more effective approach is using the light weighted selective retransmission mechanism whereby the worker only requests the retransmission of particular missing fragment indices instead of requesting the retransmission of all fragments for that entire tile. The other option is to use a time division multiple access technique that involves transmitting from the CAM to each worker in a certain time slot without any collisions at all.

Scaling the node count. Currently, the architecture uses two worker nodes and one aggregator. Using four or more worker nodes to handle two or three tiles each would be sufficient for processing images of greater resolution such as 448x448 using 16 tiles but without increasing the number of tiles per node. The softmax weighted average aggregation by the aggregator is independent of the number of nodes and requires only *NUM_TILES* constant.

Multi class detection. The binary human/no human classifier could be extended to a multi class output predicting additional categories such as vehicle, animal, or background, making the system applicable to broader surveillance and wildlife monitoring use cases with minimal architectural change.

Privacy preserving deployment. Because all processing occurs on device with no image data leaving the local ESP-NOW network, the architecture is inherently privacy preserving relative to cloud based alternatives. Future work could formalise this property and explore the system's suitability for GDPR compliant deployments in public spaces where raw video transmission is legally restricted.

6.4. Final Remarks

This project shows that the boundary between what is considered feasible on microcontroller class hardware and the domain of systems exclusively accelerated with GPU is not a hard line, it is an engineering problem. In this work we show that resolution preserving HPD can be achieved without centralised compute, cloud connectivity or specialised neural processing hardware by decomposing a spatial inference task into tile level subtasks, distributing them across a network of sub five dollar microcontrollers and aggregating their outputs by means of an attention weighted fusion mechanism. The deployed system is a research prototype with real performance limitations, most notably in inference latency and model accuracy. However, the infrastructure built the end to end quantization pipeline, the ESP-NOW distributed communication layer, the fragment reassembly mechanism and the attention based aggregator constitutes a reusable foundation on which the identified improvements can be systematically applied. As MCU clock speeds increase and TFLite Micro's hardware acceleration support matures, the performance gap identified here will diminish. The architectural approach itself, spatially distributed, attention guided, feature level aggregation on heterogeneous edge nodes is not hardware specific and scales naturally to more capable devices as they become available at comparable cost and power envelopes. The broader significance of this work lies in its demonstration that intelligent sensor systems need not be centralised to be capable. In environments where privacy, latency, connectivity, and power are all constrained simultaneously, distributed edge inference of the kind demonstrated here represents a principled and practical path forward.

A. APPENDICES

A.1. PROJECT BUDGET

Table A-1. Budget Estimation

S.N.	Hardware	Units	Cost (NPR)
1	ESP32-CAM	1	Rs. 1300
2	ESP32-S3	2	Rs. 3000
3	ESP32	1	Rs. 1000
4	FTDI	1	Rs. 350
5	Data Cables	4	Rs. 400
6	I2C LCD Display	1	Rs. 1000
7	Others	-	Rs. 1000
Grand Total			Rs. 8050

A.2. PROJECT TIMELINE

A.2.1. Gantt Chart




Table A-2. Gantt Chart for Project Timeline